



SONGIFAI

Exploring the use of Covariate Specific Word Embeddings

Jonathan Magbadelo

Candidate Number: 149708

Supervisor: Dr. Julie Weeds

Submitted for the degree of Bachelors of Computer Science and Artificial
Intelligence

University of Sussex

April 2019

Statement of Originality

This report is submitted as part requirement for the degree of Computer Science and Artificial Intelligence at the University of Sussex. It is the product of my own labour except where indicated in the text. The report may be freely copied and distributed provided the source is acknowledged.

Signature:

Jonathan Magbadelo

Acknowledgements

I would like to thank my project supervisor, Dr Julie Weeds for her constant support throughout the development of this project. I would also like to thank my colleagues at Brandwatch who provided expert advice whenever I required it.

UNIVERSITY OF SUSSEX

JONATHAN MAGBADELO

SONGIFAI

EXPLORING THE USE OF COVARIATE SPECIFIC WORD EMBEDDINGS

SUMMARY

Word embedding algorithms such as GloVe are vector space models capable of providing a distributed representation of words. By utilising vast amounts of text corpora, these representations are able to encapsulate semantic and syntactic regularities along with relationships between words. Traditional word embedding algorithms tend to operate on corpus documents in solidarity, often neglecting additional covariate metadata attached to corpus documents.

CoVeR, an extension of the GloVe algorithm, jointly learns word embeddings together with a set of diagonal weight matrices, representing the affect of a particular covariate on the base embeddings.

This project explores the use of covariate specific word embeddings for both neural language modelling and text classification. Specifically, both models are applied to a possible use case: a songwriting assistant application. The main areas covered in this dissertation are:

- An overview of issues songwriters face whilst writing songs.
- An overview of related research
- An introduction to word embeddings, neural language modelling and text classification.
- Requirements analysis for SONGIFAI, the prototype application.
- A detailed account of the implementation process.
- An evaluation of CoVeR and its usage in each model.
- An evaluation of SONGIFAI
- Limitations and future work

Contents

List of Tables	vii
List of Figures	viii
1 Introduction	1
1.1 Overview	1
1.2 The Songwriters Dilemma	2
1.3 Aims and Objectives	4
2 Professional Considerations	6
2.1 Code of Conduct	6
2.2 Good Practices	6
2.3 Ethical Considerations	7
3 Related Work	8
3.1 Lyric Language Modelling	8
3.2 Lyric Genre Classification	8
4 Background	9
4.1 Word Embeddings	9
4.1.1 Previous Methods	9
4.1.2 GloVe	10
4.1.3 CoVeR	10
4.2 Language Models	11
4.2.1 Count Based Models	11
4.2.2 Neural Language Models	12
4.2.3 Text Classification	14
5 Requirements Analysis	17
5.1 Existing Solutions	17
5.1.1 MasterWriter	17
5.1.2 Rhymer’s Block	18
5.1.3 Requirements analysis	18
5.2 Requirements	19
5.2.1 Functional	19
5.2.2 Non-Functional	19
6 Methodology	21
6.1 Collecting Data	21
6.2 Data Analysis and Restructuring	21
6.3 Data Pre-processing	22
6.4 Hyperparameters	23

6.4.1	CoVeR Hyperparamters	23
6.4.2	LSTM Hyperparameters	24
7	Implementation	25
7.1	Hardware Specification	25
7.2	Calculating Co-occurrence Statistics	25
7.3	CoVeR Implementation	26
7.4	LSTM Implementations	27
7.4.1	Language Model	27
7.4.2	Text Classifier	27
7.5	SONGIFAI	27
7.5.1	Architecture	27
7.5.2	Class Overview	27
8	Evaluation	28
8.1	CoVeR Evaluation	28
8.1.1	Validating Implementation	28
8.2	Model Evaluations	28
8.2.1	Language Model	28
8.2.2	Text Classification	28
8.3	SONGIFAI	28
8.3.1	Requirements Evaluation	28
8.3.2	Expert User Testing	28
9	Conclusion	29
9.1	What was I right about?	29
9.1.1	Previous theories were wrong	29
9.1.2	My new idea is right	29
	Bibliography	30
A	Code	32

List of Tables

3.1	Example slang words found in the dataset used in the project	8
5.1	MasterWriter Features	18
5.2	Rhymer's Block	19
5.3	SONGIFAI Functional Requirements	20
5.4	SONGIFAI Non-Functional Requirements	20
7.1	Hardware specification for machine used throughout development	25

List of Figures

1.1	Average Number of Billboard 100 songs vs Unique Word Count	3
1.2	Distribution of Music Genres throughout History	4
1.3	SONGIFAI: High Level Architecture	5
4.1	Matrix factorization	9
4.2	Skip-gram and CBOW models	10
4.3	Example neural network, three input nodes, four hidden and two outputs .	13
4.4	An <i>unrolled</i> recurrent neural network can be seen as a feed-forward neural network with many hidden layers	14
4.5	LSTM memory cell, with forget, input and output gates	15
5.1	MasterWriter	17
5.2	Rhymer's Block	18
6.1	Average word count per lyric per genre in the dataset	22
7.1	High level view of the Spark Architecture. The spark context is where the main program is defined, which is then split into tasks to be completed via numerous executors.	26

*The skill of writing is to
create a context in which
other people can think*

EDWIN SCHLOSSBERG

Chapter 1

Introduction

1.1 Overview

Natural language is a constantly evolving system capable of producing infinite numbers of sentences each with their own distinct meaning. The task of inferring meaning from a sentence is one that is seemingly trivial for humans, even from a young age. However, getting machines to derive meaning from textual data has long been a challenge within the field of Artificial Intelligence, specifically the subfield of Natural Language Processing (NLP). Computing numeric representations of words is the first step in bridging the gap between natural language and machines. Previous methods to create such representations include the one-hot encoding method which involves encoding words as word vectors which have as many dimensions as the number of unique words within a corpus. Each dimension in the vector is assigned a zero value except for the dimension representing the word which is assigned a value of 1. For example given a corpus with the following vocabulary

$$\{ \text{"Billie", "Jean", "is", "not", "my", "lover"} \}$$

and assuming that each word's index represents the dimension of the word, the word vector for "Jean" would be

$$[0, 1, 0, 0, 0, 0]$$

The discrete word representations obtained from these types of methods are sparse and suffer from the curse of dimensionality due to the dimensionality of the word vector growing linearly with the size of the vocabulary. Moreover, these representations fail to capture any syntactic and semantic relationships contained within textual data. Recent word representation methods rely on the intuition that word meaning can be inferred from context. Specifically, these methods are based on the distributional hypothesis ([Harris \(1954\)](#)) that states that words which appear in a similar context have similar meanings. Algorithms such as GloVe ([Pennington et al. \(2014\)](#)) and word2vec ([Mikolov et al. \(2013a\)](#)) exploit this concept to produce dense continuous word representations which are able to encode syntactic and semantic regularities. These types of distributed representations, first proposed by [Rumelhart et al. \(1986\)](#), are also known as *word embeddings* due to the "embedding" of words into often low dimensional space.

Often accompanying text corpora are associated covariates such as authorship and publication date, which provide additional metadata about a corpus. Traditional word embedding methods such as GloVe, fail to capture this information and instead rely on individual word embeddings for each subset of data relating to a particular covariate. This approach

to creating covariate specific word embeddings has two drawbacks: **data efficiency**, as performing these methods on subsets is repetitive, and can also yield suboptimal word representations if data is scarce for a particular subset, and decreased **interpretability**, due to the random initialisation of word embeddings, which even when created with the same random seed, produce representations relative to the the structure of the text.

CoVeR (Tian et al. (2018)), a novel tensor decomposition method for learning covariate specific word embeddings, is an extension of the GloVe algorithm which aims to efficiently encode covariate information within learned embeddings whilst increasing interpretability. The goal of this project, is to explore the effectiveness of CoVeR generated word embeddings as well as its potential application within a prototype software aimed at songwriters.

The motivation behind using CoVeR is two-fold. Firstly, its efficient use of data is beneficial when data isn't readily available, which, in the case of song lyrics, is true (discussed more in chapter 2) and its increased interpretability compared to traditional methods which allows for greater analysis. The next section briefly contextualises the project, highlighting common issues songwriters face during the writing process.

1.2 The Songwriters Dilemma

Songwriting is an integral part of the song making process which often draws upon past events and experiences. Structure and content both contribute heavily towards the success of a song; with the latter being a key factor on the extent to which a song resonates with a listener. A problem commonly faced by songwriters is that of word choice, through which they can express their ideas clearly and concisely.

Skilled writers are commonly attributed with having vast vocabulary ranges. For adults, the average vocabulary size ranges between 15,000 - 23,000 words (McCrum (2011)). Examining his works alone, Shakespeare is said to have had an approximate vocabulary range of 30,000 - 65,000 words^{1 2}. Nonetheless, a skilled songwriters ability to write impactful lyrics is not down to vocabulary size alone, but effective word choice.

As shown in a study examining vocabulary range within Hip Hop³, which recently surpassed Rock as the most popular genre in America⁴, more is not always better. The study examines the unique word count of 150 famous Hip Hop artists across their first 35,000 lyrics. Aesop Rock, ranked 1st on the list, recorded a count of 7,392 unique words across his first 35,000 lyrics. In contrast, rappers Drake and Future, who ranked 130th and 131st respectively, had an average unique word count of 3,334 words used across their first 35,000 lyrics; a 55% decrease from Aesop Rocks count.

To validate the earlier claim that vocabulary range is not indicative of a songwriters ability to write impactful lyrics, the unique word count per artist was compared against the average number of Billboard 100⁵ songs an artist had obtained across their career (seen in Figure 1.1). Pearson's Correlation Coefficient, which is used to measure the linear relationship between two variables, was applied to both unique word count and average number of Billboard 100 songs. This resulted in a correlation coefficient of - 0.42,

¹<https://kottke.org/10/04/how-many-words-did-shakespeare-know>

²This range is most likely skewed as variant forms of words are included as singular counts

³<https://pudding.cool/projects/vocabulary/index.html>

⁴<https://www.nielsen.com/us/en/insights/reports/2018/2017-music-us-year-end-report.html>

⁵<https://www.billboard.com/charts/hot-100>




Figure 1.1: Average number of Billboard 100 songs during artist activity, compared to unique word count across an artists first 35,000 words.

indicating a weak inverse correlation between the pairs of data. This value supports the earlier claim that vocabulary range is not indicative of a songwriters ability.

Common methods used to improve songwriting competency include group writing and vocabulary expansion. More recently, software solutions such as MasterWriter⁶ have been used to consolidate previous methods. An inherent drawback with applications such as MasterWriter is the static nature of features such as fixed word dictionaries and thesauruses. Consequently, these applications fail to capture the ever changing nature of language.

After the completion of song lyrics, another task in the song making process, if not already done, is song production. For experienced songwriters, the choice of production style might be straightforward due to them already having an established production style. On the other hand, for newer songwriters or songwriter's wishing to be heard over a new sound, this choice may not be a trivial one, especially with the constant increase in sub genres, and their relevant production styles, within music. As shown in Figure 1.2, the distribution of music genres listened to has spread over time, even more so, the increase of distinctive sub genres within these genres.

(QUANTIFY SUBGENRE TALK)

⁶<https://masterwriter.com/5>




Figure 1.2: Distribution of music genres listened to throughout history based on the music library of Google Play Users

7

1.3 Aims and Objectives

As previously stated, the goal of this project is to explore effectiveness of CoVeR as well as its potential usage in a software application. Specifically this takes the form of a prototype software solution to help reduce common problems faced by songwriters, namely:

1. Word Choice
2. Lyric classification

The tasks of choosing the next word in a song lyric and lyric classification can be viewed as abstracted forms of two common NLP tasks: *Language Modelling* and *Text Classification*, which are two active research areas in NLP that have benefited greatly from word embeddings. The usage of these word embeddings has helped improve accuracy in tasks such as Document Classification and Machine Translation. State of the art results in both areas have utilised recurrent neural networks, (discussed in [page 9](#)) which utilise word embeddings within the input layer to the network.

With this in mind, a prototype solution, SONGIFAI is proposed. SONGIFAI provides two main features namely lyric assistance through predictive text and a word suggestion feature, as well as genre classification for a given lyric. The covariates explored in this project are the following music genres: *Pop*, *Rock* and *Hip Hop*.

As stated before, recurrent neural networks have been used successfully in both language modelling and text classification tasks. The specific variant of the recurrent neural network which has been successfully utilised in these tasks is the Long Short-Term Memory network. This network architecture is used for both the language model and text classifier which are utilised for the word prediction and lyric classification features respectively. The word suggestion feature utilises the covariate specific embeddings themselves. A high level system architecture of the prototype can be seen in [Figure 1.35](#)

Training

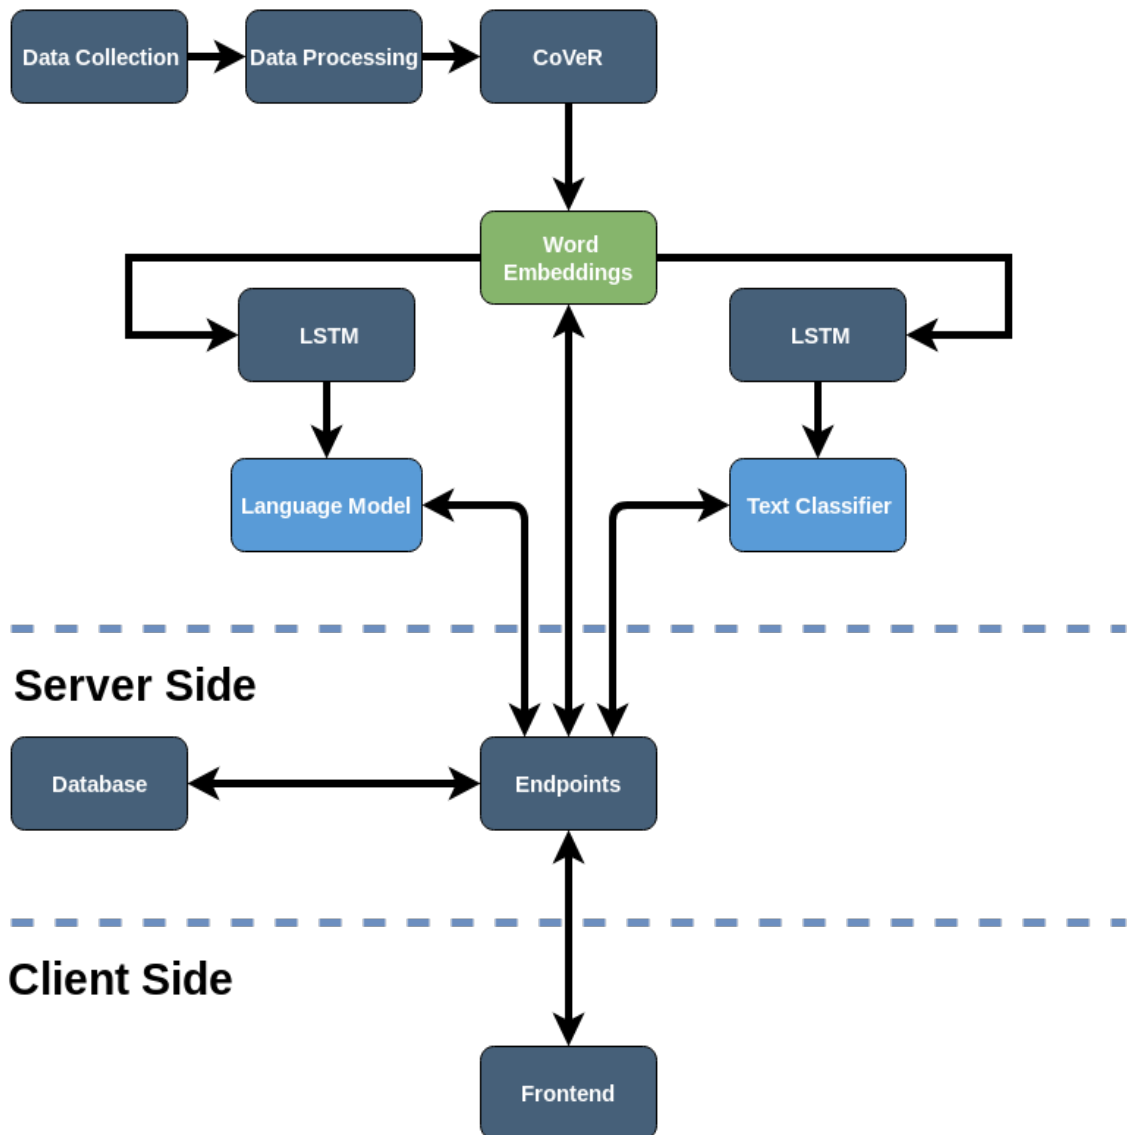


Figure 1.3: From a high level, the system can be split into three sections: **Training**, **Server Side** and **Client Side**

Chapter 2

Professional Considerations

Throughout the development of this project both professional and ethical considerations were taken into account, including those highlighted in the British Computing Societies (BCS) Code of Conduct ¹ and Code of Good Practice. This chapter outlines the relevant areas in the specified documents which have been adhered to during the project.

2.1 Code of Conduct

Professional Competence and Integrity

The completion of this project was a large undertaking due to the implementation and integration of a novel machine learning method within a prototype software application. Though the project is beyond the scope of a typical final year project, all work carried out have roots to modules taken in the University of Sussex Computer Science and Artificial Intelligence course, specifically the Neural Networks, Advance Natural Language Engineering and Software Engineering modules. In accordance with section 2.C of the BCS Code of Conduct, background research continually occurred throughout development to maintain a competent standard of professional knowledge.

Duty of Relevant Authority

In agreement with section 3.A and 3.B of the BCS Code of Conduct, all scenarios which may cause a conflict of interest between the project and the University of Sussex have been avoided.

Duty to Professionalism

In accordance with section 4.A and 4.C of the BCS Code of Conduct, the manner in which this project was conducted was one which maintained the reputation of the BCS. During meetings with other BCS members and professionals such as my project supervisor and work colleagues, appropriate levels of respect and integrity were upheld in accordance with section 4.B of the BCS Code of Conduct.

2.2 Good Practices

The motivation behind this project is one rooted in exploratory research rather than being client driven. Nevertheless, it is important that code produced is well structured and testable to ensure quality assurance. Where possible and in accordance with section

¹<https://www.bcs.org/category/6030>

5.2 of the BCS Code of Good Practice, the code produced is well structured and organised to help facilitate further testing and maintainability.

The same section of the BCS Code of Good Practice refers to the following of programming language guidelines. Both Python and JavaScript were used extensively during the development of the project and where appropriate best practices and coding style/conventions have been adhered to.

2.3 Ethical Considerations

The success of a machine learning project relies heavily on data availability and quality. Currently, there exists no central repository for downloadable song lyrics due to potential copyright infringement. Under the Research and Private Study section of the Copyright, Designs and Patents Act (1988) ², a public song lyric dataset is used solely for training machine learning models.

Finally, the project utilises textual data which may have explicit or offensive content within it. After the training of models, which will not be filtered to allow permit artistic freedom, a filtering option will be implemented in order to prevent potential users from seeing unwanted content.

²<https://www.gov.uk/government/publications/copyright-acts-and-related-laws>

Chapter 3

Related Work

This chapter briefly outlines previous work relating to the tasks of language modelling and text classification using song lyrics.

3.1 Lyric Language Modelling

A general approach to modelling language used in lyrics includes the usage of Markov processes. Naturally, as only the present or past is considered, these types of models are unable to capture longer dependencies which can occur in lyrics. Recent attempts to model creative writing have utilised neural methods, such as [Zhang and Lapata \(2014\)](#), who uses RNN's for Chinese poetry generation and [Potash et al. \(2015\)](#), who uses an LSTM, to ghostwrite¹ rap lyrics. Specific to rap lyrics [Hulzebosch \(2017\)](#) extends previous works by restructuring the problem as a sequence to sequence model which is aware of generated sequences whilst ... As stated in the paper, a drawback of this approach is the large amounts of training data required to produce good results.

3.2 Lyric Genre Classification

Previous attempts at genre classification predominately use audio features rather than lyrics. Apart from audio features, the genre of a song is closely related to the lyrics as songs of varying genres use language differently. This can be seen in the Hip Hop, where the use of slang, which are twisted existing words or self-made words ([Edwards and Rap \(2009\)](#)) is prevalent. Examples of such words can be seen in [Table 3.1](#)

Table 3.1: Example slang words found in the dataset used in the project

Slang Word	Meaning
grill	teeth
lit	fun
bread	money

Similar to the previous section, genre classification has also been tackled using neural methods. Using audio based features, ([Irvin et al. \(2016\)](#)) and ([Chun Pui Tang \(2018\)](#)), utilise LSTM's for classification.

¹write in the style of X without giving credit

Chapter 4

Background

This chapter provides an introduction to the theory and previous work within the areas of word embeddings, neural language models and text classification.

4.1 Word Embeddings

Word embeddings are vectors of predefined size which aim to encode a distributional numerical representation of word features. Recent aforementioned methods of learning these representations include the GloVe, word2Vec and fastText (?) algorithms. The utilization of word embeddings has been highly successful in many NLP tasks such as sentiment analysis (Socher et al. (2013)) and syntactic parsing (Socher et al. (2013)). Previous techniques for creating such representations can be categorised into two categories: matrix factorization methods and shallow window-based methods.

4.1.1 Previous Methods

Global Matrix Factorization Methods

Global matrix factorization methods such as Latent Semantic Analysis (LSA) use low rank approximations to decompose large matrices containing corpus statistics. Typically these matrices take the form of a term-document matrix, which captures the frequencies of terms across a collection of documents, or a term-term matrix, which store co-occurrence counts of terms. Matrix factorisation methods such as LSA allow for fast training and perform well on word similarity tasks by leveraging word occurrence statistics however they suffer from the disproportionate importance given to large word counts.

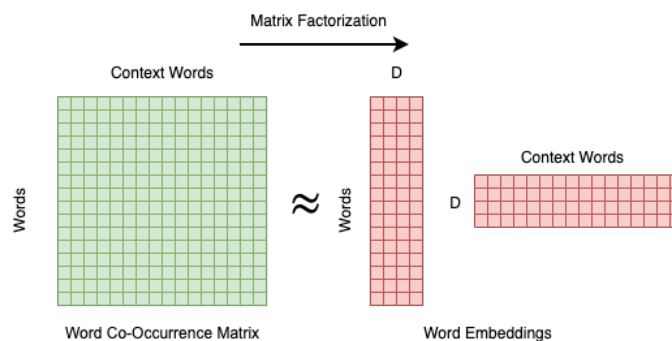


Figure 4.1: Matrix factorization

Shallow Window-Based Methods

Shallow window-based methods provide an alternative approach to learning word representations by sliding a fixed window over the contents of a corpus and learning to predict either the surroundings of a given word (skip-gram model) or predict a word given its surroundings (continuous bag of words). In the case of shallow window-based methods, they are good at capturing more complex patterns and do well in the word analogy task, however they fail to leverage global statistical information such as those used in global matrix factorization methods.

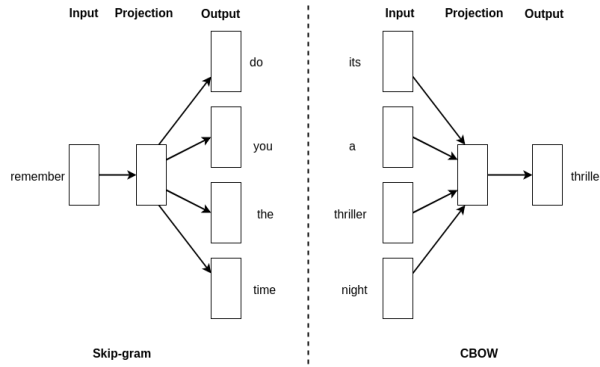


Figure 4.2: Skip-gram and CBOW models

4.1.2 GloVe

Global Vectors for Word Representation (GloVe), is an unsupervised word embedding algorithm, introduced by [Pennington et al. \(2014\)](#), which marries the benefits of both global matrix factorisation and shallow window-based methods. Presented as a log-bilinear regression model, GloVe makes use of global word co-occurrence statistics from a corpus. Specifically, given a corpus has vocabulary size N , a word-word co-occurrence matrix X is computed with X_{ij} being a measure of the number of times words i and j co-occur within a given context window, which in GloVe is a inverse distance function, giving closer words higher counts.

As detailed in the paper, GloVe outperformed previous methods such as word2vec in word analogy, word similarity and named entity recognition tasks.

Conceptually, GloVe is based on the idea that ratios of probabilities of words co-occurring have the potential to encode meaning which is encoded as vector differences. (REFACTOR)

This concept is formalised in the following equation, where the dot product of focal and context word vectors, w and \tilde{w} , is equal to the logarithm of the number of times the words co-occurred, $\log X_{ij}$.

$$w_i^T \tilde{w}_j + b_i + \tilde{b}_j = \log X_{ij}^2 \quad (4.1)$$

GloVe also utilises a weighting function $f(X_{ij})$ to decrease noise caused by high frequency word co-occurrences. The following weighting function is used in the GloVe model.

$$f(x) = \begin{cases} (x/x_{max})^\alpha, & \text{if } x < x_{max} \\ 1, & \text{otherwise} \end{cases} \quad (4.2)$$

Where α is 0.75 and x_{max} is 100. Combining equations 3.1 and 3.2, the GloVe model is defined as a weighted least squares regression problem.

$$J = \sum_{i,j=1}^N f(X_{ij})(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2 \quad (4.3)$$

4.1.3 CoVeR

Covariates such as author demographics, time and location often accompany documents within a corpus. A trivial approach to obtaining covariate specific word embeddings involves applying GloVe to each subset of corpus documents relating to a particular covariate. Unfortunately, utilising GloVe this way has certain drawbacks. Firstly, GloVe must be applied and optimised to each sub corpus, which, depending on the number of covariates, is computationally expensive. Also as a direct result of dividing the original corpus, global co-occurrence statistics are now split across covariate specific co-occurrence statistics. Consequently, global statistics are now not shared between embeddings, which may cause sub-optimal word representations, especially when sub corpus contains small amounts of co-occurrences for GloVe to train on.

A known problem when representing words as dense representations using methods such as GloVe, is interpretability of dimensions. Due to conditional GloVe, producing word embeddings per corpora, relating these embeddings to one another becomes a difficult task.

Learning Covariate-Specific Vector Representations with Tensor Decompositions (CoVeR), proposed by Tian et al, provides an alternative to the conditional GloVe method which offers a framework to make learned embeddings more interpretable. Being an extension of GloVe, CoVer extends GloVe’s matrix decomposition of co-occurrence matrices to tensor decomposition of co-occurrence tensors, involving the joint learning of word embeddings and covariate specific transformation matrices which represent the effect of a particular covariate on the base embeddings learned. The CoVeR model is presented below.

$$J = \sum_{i,j=1}^N \sum_{k=1}^M f(X_{ijk})((c_k \odot w_i)^T (c_k \odot \tilde{w}) + b_{ik} + \tilde{b}_{jk} - \log X_{ijk})^2 \quad (4.4)$$

The introduction of covariate specific weight matrices into the objective function allows the authors to interpret dimensions in the base embeddings learned.

4.2 Language Models

Formal languages such as programming languages are fully specified with precise syntax and semantics which dictate the usage of all reserved words within a language. Contrarily, natural languages, because of their emerging nature, are unable to be formally specified even with the existence of grammatical rules and structures. Unfortunately, rule based systems suffer from the endless possibilities of language usage outside of grammatical rules which are easily interpretable by humans. Moreover the task of consistently updating rule based systems to accommodate such usage is unfavourable.

Language modelling (LM) is the task of estimating the probability distribution of various linguistic units such as characters, words and sentences. In recent years, the application of

LM has been essential to many natural language process tasks such as speech to text and text summarization. Language models can be classified into two categories, count-based and continuous-space language models.

4.2.1 Count Based Models

Count based methods such as statistical language models attempt to learn a probability distribution $P(w_1, \dots, w_i)$ over of a sequence of words w_1, \dots, w_i . An example of a count based method is the n-gram model.

An n-gram is a sequence of n words. Examples of a two word sequences or bigrams include, "My name" and "is Aubrey", whilst examples of three word sequences or trigrams, include sequences of words such as "Hello my name" and "is Aubrey Graham". The n-gram model which considers the past $n - 1$ words can be formalised as

$$P(w_i|w_1, \dots, w_{i-1}) \approx P(w_i|w_{i-n+1}, \dots, w_{i-1}) \quad (4.5)$$

The n-gram model relies on Markov assumptions to model the probability of word sequences $P(w_1 \dots w_n)$ as being equal to a limited number of previous words. An inherent problem with the n-gram model is sparsity as some word sequences occur rarely or not at all, even in large text corpora. Using the standard n-gram model would yield too many zero probabilities. To circumvent this, techniques such as back-off and smoothing exist. Another disadvantage of n-gram models is that they rely on exact patterns, meaning they fail to recognise syntactically and semantically similar sentences such as "the cat sat on the mat" and "the dog sat on the mat". N-gram models also suffer from the curse of dimensionality due to increased vocabulary sizes. As a result, limited window sizes are used, causing longer dependencies between words to not be captured.

4.2.2 Neural Language Models

To overcome issues faced by count based models, deep learning methods have been used to create neural language models by simultaneously learning word embeddings and the parameters needed to create a joint probability distribution between the word embeddings. [Bengio et al. \(2003\)](#) proposed a feed forward neural language model to help tackle the problem of data sparsity. Recent state of the art approaches such as [Mikolov et al. \(2010\)](#), abstract language modelling as a form of sequential data prediction and have implored recurrent neural networks to help encode longer dependencies between sequences of words. The strength of these models comes from their ability to consider several preceding words and thus generalise well.

An overview of generic neural network architectures as well as recurrent neural networks and Long-Short Memory networks are given in the following sections.

Artificial Neural Networks

In any neural network architecture, the elementary unit of computation is the artificial neuron which takes inspiration from biological neurons. The artificial neuron receives n inputs which are each weighted by n weights and summed together with a bias b . The output y of a neuron is calculated by passing the weighted sum of the inputs into an activation function f .

$$y = \left(\sum_{i=1}^N x_i w_i + b \right) \quad (4.6)$$

Typical activation functions include *Sigmoid*, *Tanh* and *ReLU*. A single layer neural network is defined by k neurons sharing the same input in the same layer. Single layer neural networks have been proven to be '*universal approximators*' (?) meaning any continuous function can be approximated using this type of network. The process of stacking layers on top of each other leads to multi-layer neural networks. These types of networks are also known as feed-forward networks. The learnable parameters of these networks are the set of weights and biases for each layer. A feed-forward neural network is trained using gradient descent and its parameters are updated using the *backpropagation* algorithm (?).




Figure 4.3: Example neural network, three input nodes, four hidden and two outputs

Recurrent Neural Network

In a feed-forward neural network, data flow is unidirectional between layers; with data passing through a given neuron at most once. These types of networks perform well on both classification and regression tasks with the assumption that inputs are independent of each other. In tasks dealing with sequential data, feed-forward networks perform poorly. To model sequential data well, a neural network must be able to model the dependencies that exist between successive inputs. The recurrent neural network (RNN) is an attempt to satisfy this requirement by utilising past inputs to help predict future outputs.

In an RNN information is cycled within the network at least once. An RNN receives a sequence of inputs x and updates its hidden state h_t by

$$h_t = \begin{cases} 0, & t = 0 \\ \phi(h_{t-1}, x_t), & \text{otherwise} \end{cases} \quad (4.7)$$

where ϕ is a nonlinear function such as *tanh* or *ReLU*. The update for the hidden state is usually implemented as

$$h_t = \phi(Ux_t + Wh_{t-1}) \quad (4.8)$$

where W and U are weight matrices.

RNN's are trained using gradient descent and backpropagation through time (BBTT), which is identical to performing backpropagation on an "unrolled" RNN (seen in figure [Figure 4.4](#))




Figure 4.4: An *unrolled* recurrent neural network can be seen as a feed-forward neural network with many hidden layers

During BBTT, back propagation is performed on an unrolled recurrent architecture, causing gradients to back-propagate through numerous network layers. Unfortunately, this has a few major problems. Firstly, a single forward/backward pass through the network is computationally expensive due to the number of hidden layers of the unrolled network being linear to the number of time steps in a sequence. Secondly, this method suffers from the issue of vanishing or exploding gradients(REF) where gradients can decay or grow exponentially as they propagate over time, which can prevent the network from learning entirely.

Long Short-Term Memory

Long Short-Term Memory (LSTM) ([Hochreiter and Schmidhuber \(1997\)](#)) is a variant of the recurrent neural network which is capable of capturing longer dependencies between sequences of data without suffering from vanishing gradients. This is achieved through a feature known as gating; a mechanism which acts as a permissive or restrictive barrier to information flow.

The core component of the LSTM is the cell state which is able to propagate **relevant** information throughout the network. This is achieved within the memory cell through the forget, input and output gate. The forget gate regulates how much of the existing memory should be forgotten, the input gate regulates how much of the new cell state to keep, and the output gate regulates how much of the cell state should be allowed into the next layer of the network.




Figure 4.5: LSTM memory cell, with forget, input and output gates

4.2.3 Text Classification

Text classification is the task of assigning pre-defined labels to text according to its contents. Machine learning methods for text classification are preferred to rule based systems due to their ability to model regardless of domain. Text classification can be approached using supervised learning, where a model is trained with labelled training data, and unsupervised learning, where training data are clustered together to see if any natural grouping exists.

Before a classifier can be trained, textual inputs must be transformed into numerical representations. A common method for creating such representations is the bag of words approach which given a vocabulary, creates an input vector which represents word counts for each word of the vocabulary. For example, if we define the vocabulary V as being the set of words $\{ "I", "like", "food", "had", "today" \}$ and a sentence S as being *"I had food"*, the resulting vector for word S would be $[1, 0, 1, 1, 0]$. Two popular machine learning methods for text classification include the naive Bayes classifier and the Support Vector Machine.

Naive Bayes

Naive Bayes classifiers have been successfully used for text classification tasks such as spam filtering (REF) and provide a way to model data without the need of vast amounts of training data. They are a class of generative classifiers which rely on Bayes theorem to learn a model of the joint probability, $P(x, y)$, of the inputs x and target classes y to calculate the conditional probability $P(y|x)$. Concretely, given a set of features x_0, \dots, x_n naive Bayes classifiers assume conditional independence between features and use prior class probabilities along with ... to assign a class label to a new set of features. The general formula for naive Bayes classifiers is presented below.

$$y = \operatorname{argmax}_y P(y) \prod_{i=1}^n P(x_i|y) \quad (4.9)$$

Support Vector Machines

Support Vector Machines (SVM) ([Corinna Cortes \(1995\)](#)) have also been applied successfully to text classification problems ([Joachims \(1998\)](#)). Unlike naive Bayes classifiers, SVM's are discriminative classifiers which aim to find a hyperplane in N-dimensional space which maximises the margin between data points of different classes. This is achieved through the use of random data points, namely "*support vectors*", which are data points closest to the hyperplane which help maximise the margin distance between classes. Maximising marginal distance between classes provides confidence to the classification of future data points.

IMAGE OF SVM

Deep Learning Methods

Until recently, state of the art methods for text classification had long used linear predictors such as the SVM (with a linear kernel) and multinomial naive Bayes classifiers using feature representation techniques such as bag-of-words ([Joachims \(1998\)](#) [Lewis et al. \(2004\)](#)). However, non-linear methods that are able to leverage dependencies between words, such as recurrent neural networks, have provided better predictions ([Dai and Le \(2015\)](#)) than previous bag-of-word approaches where features are assumed to be conditionally independent and unordered.

Chapter 5

Requirements Analysis

As the project involves the development of a prototype software system, it is important to consider the project from a software engineering perspective. Moreover, the project involves the integration of a novel machine learning model, with the success of the project relying heavily on factors such as data availability, data quality and processing power. With this and other considerations such as training time and implementation complexity in mind, it is necessary to define software requirements to limit the project scope to one that is achievable. Software requirements should also be inferred from the needs of the end-user and as such it is necessary to understand user needs through existing solutions. This chapter briefly evaluates two existing solutions and outlines the functional and non-functional requirements by which the prototype will be evaluated.

5.1 Existing Solutions

5.1.1 MasterWriter

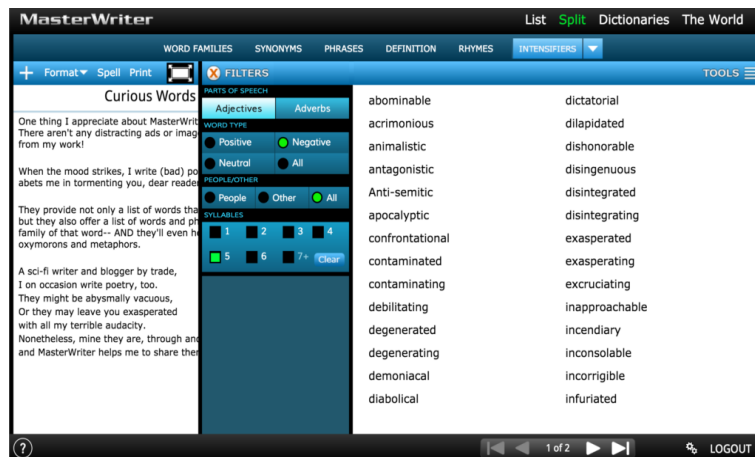


Figure 5.1: MasterWriter

Self-described as *"The most powerful suite of writing tools ever assembled in one program."*, MasterWriter is a software application aimed towards songwriters, poets and creative writers. Available as a multi-platform product, it consolidates a number of writing aids into one application. Some of the core features of MasterWriter are highlighted in the table below.

Table 5.1: MasterWriter Features

Feature	Description
Word Dictionary	Provides a word dictionary with word definitions
Rhyming Dictionary	Provides a rhyming dictionary with word/phrase rhymes
Phrase Dictionary	Provides a phrase dictionary of predefined phrases
Word Families	An extension of a typical thesaurus which provides a reference dictionary which can filter on subtle differences in word meanings
Speech Types	Provides a predefined list of different figures of speech
Synonyms	Provides a thesaurus containing synonyms

5.1.2 Rhymer's Block




Figure 5.2: Rhymer's Block

Rhymer's block is a mobile application intended to help songwriters specifically with rhymes. Providing real time rhyme suggestions, the application allows users to quickly write song lyrics and provides a social platform through which users can share their own song lyrics, as well as review the song lyrics of other users. An overview of Rhymer's Block core features can be seen in the table below.

5.1.3 Requirements analysis

Both MasterWriter and Rhymers Block provide functionality for users to write, edit and save their lyrics. Similarly, both applications offer word search functionality to help writers with word choice, a primary issue highlighted before, through features such as rhyme suggestions, word dictionaries and thesauruses. Taking inspiration from both applications and as a basis for SONGIFAI, a lyric editor and word suggestion feature are mandatory for implementation. As previously stated, software applications like MasterWriter are static in nature, and fail to capture the dynamic nature of language. Examining Rhymer's

Table 5.2: Rhmyer’s Block

Feature	Description
Real-Time Rhyme Suggestion	Provides real time suggestions of rhyming words
Cloud Storage	Provides cloud storage for a user’s song lyrics
Social Platform	Provides a online platform for users to share and comment on song lyrics
Word Dictionary	Provides a word dictionary with word definitions
Thesaurus	Provides a thesaurus containing synonyms
Word Analysis	Provides rhyme suggestions for words most commonly used by other users

Block, a particular feature deviates away from this, namely its real-time rhyme suggestion feature which, leverages word usage within the applications online community. The implementation of a social platform is not in the scope of this project, however Rhmyer’s Block leverage of its users language is comparable to ...

Neither MasterWriter or Rhmyer’s block address the secondary issue raised in [chapter 1](#); namely helping songwriters with classifying their lyrics to help infer instrumental style. Whilst this project only considers three music genres, the potential application of such a feature becomes more apparent when considering sub-genres within music. (EXPLAIN MORE HERE) Though sub genres are out of scope for this project, a basic genre classification feature will be implemented for SONGIFAI which will classify a user’s submitted song lyric as either Pop, Rock or Hip Hop.

Comparing MasterWriter and Rhmyer’s block, software availability and accessibility are dealt with differently. Whilst MasterWriter offers its application on various platforms, Rhmyer’s Block utilises cloud storage for better accessibility. To address both software availability and accessibility, whilst keeping the project in scope, SONGIFAI will take the form of web application.

As a whole, the project is both research based and practical in nature. Often times the methodologies of research and practice collide, as whilst research tends to focus more on ...,

5.2 Requirements

In this section the requirements for the project will be set out. The functional requirements will specify what the software will do whilst the non-functional requirements will detail how these will be done.

5.2.1 Functional

5.2.2 Non-Functional

Table 5.3: SONGIFAI Functional Requirements

ID	Description	Dependency
FR1	The system should allow users to input lyrics	N/A
FR2	The system should allow users to edit, save and delete lyrics	N/A
FR3	The system should be able to classify user submitted lyrics as either Pop/Rock/Hip Hop	N/A
FR4	The system should be able to suggest words from a given word. These words should be the most similar words in the chosen covariate word embedding space	N/A
FR5	The system should be able to provide real time text prediction whilst a user is in edit mode	N/A
FR6	The system should allow for the filtering of explicit content in both the word suggestion and word prediction feature	N/A
FR7	The user should be able to change the underlying covariate specific word embeddings used or the base embeddings if desired	N/A

Table 5.4: SONGIFAI Non-Functional Requirements

ID	Description	Dependency
NFR1	The system should take the form of a web application, compatible on multiple device types	N/A
NFR2	The word prediction feature should return a list of candidate words with minimal latency	N/A
NFR3	The word suggestion feature should return a list of suggested words with minimal latency	N/A

Chapter 6

Methodology

This chapter describes the methodology used during this project.

6.1 Collecting Data

As previously mentioned, a central repository for song lyrics, where downloading is permissive, is non-existent. Collection of song lyrics from sites such as Genius¹ is only achievable through web scraping: a process which involves the exhaustive downloading and processing of web pages from either a predefined list of URLs or the recursive process of link extraction and following (commonly known as *web crawling*) from a given seed URL. When scraping at scale, restrictions such as the robots exclusion protocol, which specifies areas of a website that are allowed to be processed, and crawl rate, which specifies the minimum delay between requests, make data collection a troublesome task. Taking into account the project aims and constraints such as time, web scraping was deemed unfavourable and hence avoided for data collection.

To fulfil the data requirements of the project, a public dataset ² containing over 250,000 song lyrics was used. The dataset, which was scraped from Brazilian music website Vagalume, comprised of two *.csv* files: *artists-data* and *lyrics-data*. The *artist-data* file included metrics such as popularity, genre and number of songs per artist, whilst the *lyrics-data* file contained song lyrics for individual songs by artists.

6.2 Data Analysis and Restructuring

The CoVeR algorithm requires each document within a corpus to be associated with a covariate in order to jointly learn word embeddings and the relevant transformation matrices. To satisfy this requirement, Pandas, a Python data analysis library, was used to perform a join operation on the artists columns in both the *artists-data* and *lyrics-data* files. The result of this operation followed by the additional dropping of redundant columns resulted in the dataset used throughout the rest of the project.

(IMAGE OF DATASET SCHEMA)

A trivial approach to split the lyric corpora on genre would involve an equal split for equal representation, however, this method uses the assumption that for each genre, word count per song is equivalent. Examining the dataset however, proved this not to be the case.

¹<https://genius.com/>

²<https://www.kaggle.com/neisse/scrapped-lyrics-from-6-genres>

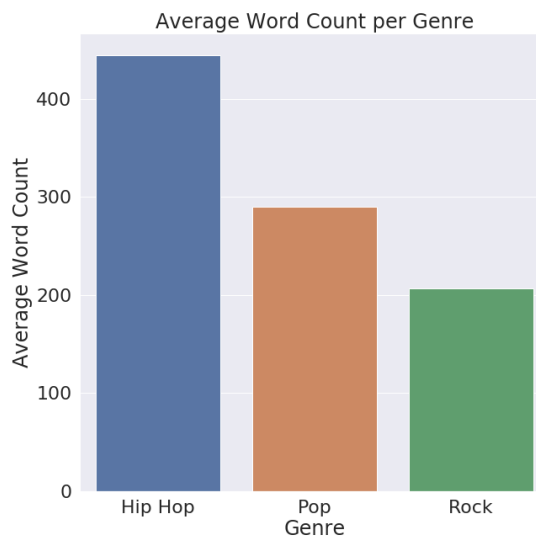


Figure 6.1: Average word count per lyric per genre in the dataset

Analysing the dataset, the mean number of words contained in a hip hop song was 444; being more than double the mean number of words found in rock songs which had a count of 207. With regards to pop, the mean word count per song was 289. To reflect these statistics and also due to constraints in the distribution of data by genre, 100,000 song lyrics were randomly selected to create the dataset to be used to train CoVeR. Moreover a data split of 48:30:22 for Rock, Pop and Hip-Hop, was to maintain an even distribution of words by song genre within the dataset.

6.3 Data Pre-processing

Essential to any machine learning task is the pre-processing of input data in such a way that important features are accessible during training. In natural language processing, this can include techniques such as tokenisation, string cleaning, stemming and lemmatisation.

Following data reconstruction, each song lyric in the dataset was cleansed and tokenised. The following string cleaning techniques were applied to each song lyric.

1. All letters were lowercased.
2. All characters, except for letters, were substituted with a space.
3. All text between brackets were removed. (This was to ensure text like **[Verse 1]** was not included during training).
4. All trailing white space was removed.

Tokenisation is the process of separating textual inputs into meaningful chunks called *tokens*. Naturally to create word embeddings, text is tokenised at the word level and each token is assigned a unique integer key to map the token to its future word embedding.

6.4 Hyperparameters

In machine learning hyperparameters are parameters that govern a given model. The selection of these parameters directly impacts the performance of a given training algorithm and as such, hyperparameter choice is important for producing optimal performance for a given model. Both CoVeR and LSTM networks have important hyperparameters which are reviewed in the following sections.

6.4.1 CoVeR Hyperparamters

Removing Rare Words

Word embedding methods such as word2vec typically remove rare occurring words within a corpus before training takes place. In the case of word2vec, rare words are removed, *before* the computation of co-occurrence statistics. This method of removal has the effect of generating context windows differing to the context windows generated from the original corpus. In turn this produces differing i, j co-occurrences. This is highlighted in the example below.

Sentence: Fly me to the moon.

Context of "to" (without removal of rare words): me, the

Context of "to" (with removal of rare words): Fly, moon

Neither the CoVeR or GloVe paper, discuss the removal of rare words, however many implementations (REF)

Subsampling

Typically found within text corpora are high frequency stop words which provide less information than rarely occurring words (Mikolov et al. (2013b)). For example This concept can also be applied to word embeddings; where the word embeddings of frequent words does not change significantly after training on several examples. Taking inspiration from word2vec, subsampling was used at the covariate level using the following adapted formula from the original word2vec implementation.

$$P(w_{ik}) = \sqrt{\frac{z(w_{ik})}{t}} + 1 \cdot \frac{t}{z(w_{ik})} \quad (6.1)$$

where $z(w_{ik})$ is the percentage of word w_{ik} in covariate k and t is a chosen threshold.

Context Windows

CoVeR uses the same weighted context windows strategy as GloVe during the process of calculating co-occurrence statistics. The CoVeR paper uses a context window size of 8, however the paper does not specify whether they used symmetric or asymmetric windows during their experiments. As such, both are experimented with whilst performing hyperparameter tuning.

Embedding Size

Optimiser

Tuning

6.4.2 LSTM Hyperparameters

Dropout

Dropout is a regularization technique proposed by [Srivastava et al. \(2014\)](#), which involves the random dropping of nodes and their connections within a network. Selected nodes are picked at random using a probability known as the dropout rate. Dropout helps to decrease over-fitting as

Embedding Layer

Optimiser

Chapter 7

Implementation

7.1 Hardware Specification

All implementation was completed on a personal machine. The hardware specifications for the machine are highlighted below

Hardware Component	Specification
CPU	Intel Core i7-8750H CPU @ 2.20GHz x 12
GPU	NVIDIA GeForce GTX 1050 Ti 4GB
RAM	16GB

Table 7.1: Hardware specification for machine used throughout development

7.2 Calculating Co-occurrence Statistics

Many unsupervised natural language processing methods compute co-occurrence statistics before learning takes place. Typical co-occurrence statistics, such as GloVe’s word-word co-occurrence matrix contain sparse data, and computing them can often be a computationally more expensive task than the learning itself.

The original GloVe paper describes this process as a *‘one-time upfront cost’*, with the assumption that selected corpora are static. Unfortunately, for many natural language processing pipelines such corpora are more dynamic in nature. For example, social data from online platforms such as Twitter are in constant flux and relying on pre-computed co-occurrence statistics is sub-optimal. Compared to GloVe, computing the co-occurrence statistics for CoVeR has added complexity due to the transition from a co-occurrence matrix to a co-occurrence tensor.

Methods for efficient computation of co-occurrence statistics include the usage of distributed computing techniques such as *MapReduce*. MapReduce is a model for distributed computing which involves two functional processes namely map and reduce. During the *map* process, data is taken in as key/value pairs and transformed to intermediary key/value pairs as output. These are then passed to the *reduce* process which aggregates data which share the same key.

Apache Spark is an open-source framework, written in Scala, for distributed computing and has recently emerged as the preferred option for big data processing over Apache Hadoop. Like Hadoop, Spark also supports the MapReduce programming paradigm but

boasts features such as enhanced speed, a distributed data structure, as well as API's written in multiple programming languages. Spark uses a master/slave architecture to achieve distributed computing. The *driver* acts as the master node and distributes tasks to many different worker nodes, also known as *executors*, which each run their own JVM processes to execute tasks.




Figure 7.1: High level view of the Spark Architecture. The spark context is where the main program is defined, which is then split into tasks to be completed via numerous executors.

PySpark, a Python API for the Spark framework was initially used to both pre-process data and calculate co-occurrence statistics for the corpora. Unfortunately, the overhead of collecting completed executor tasks to the driver as well as the cross language communication between Python and Scala made PySpark an unfavourable option for collecting co-occurrence statistics. A similar parallelised approach which avoided cross language communication involved the use of Python's multiprocessing module. This approach suffered from the restrictions of Python's Global Interpreter Lock (GIL) which prevents shared access of Python objects across multiple threads.

Cython is a superset of the Python programming language which aims to provide C like performance whilst maintaining the ability to write Python like code. Native Python programs can experience major speed improvements using Cython because of its ability to compile Python to C code. Ultimately, calculating the co-occurrence statistics for CoVeR was achieved using Cython which provided major speed ups compared to the parallelised approaches outlined before. A comparison of the speed gain for computing co-occurrences can be seen in the figure below.

(CYTHON SPEED IMPROVEMENTS DIAGRAM)

7.3 CoVeR Implementation

At time of writing, no public implementation of CoVeR is available. As a result and to meet the needs of this project, CoVeR was implemented from scratch using the PyTorch library. PyTorch is a Python library based on Torch, which supports Numpy like operations which can be accelerated through the GPU. All supporting code for the implementation can be found here: [\(LINK TO CODE\)](#)

7.4 LSTM Implementations

Implementation of both the neural language model and text classifier was done using Keras. Keras is a high level machine learning library written in Python, which runs on top of either Tensorflow or Theano. The motivation behind using Keras comes from its ease of use to quickly develop deep learning networks. In this project Keras is deployed using Tensorflow as a backend, specifically for its GPU capabilities.

7.4.1 Language Model

The structure for the language model can be seen in Figure X.X. The model consisted of an input layer, an embedding layer, a bidirectional LSTM, a dropout layer and finally a dense layer.

7.4.2 Text Classifier

The structure for the text classifier can be seen in Figure X.X. The model consisted of an input layer, an embedding layer, a bidirectional LSTM, a dropout layer and finally a dense layer.

7.5 SONGIFAI

7.5.1 Architecture

Client Side

For the client-side development of SONGIFGAI, a main requirement refers to the systems availability for web and mobile access. Being a prototype solution, it was important that the development was swift and well structured so that the research goals of the project were not hindered. To help achieve this, ReactJS was chosen as the front-end development framework.

React is a Javascript framework for building user interfaces originally developed and maintained by Facebook. The main advantages of using React

Server Side

Requirements ... refer to a user of the system being able to save, load and edit their lyrics. Moreover for easy compatibility with the Keras generated models, another Python based library was preferred as the for the server side. To meet these conditions, Django was chosen as the development framework for the back-end of the system. Django is a python based web framework which follows the model-view-template (MVT) architectural pattern.

7.5.2 Class Overview

Chapter 8

Evaluation

8.1 CoVeR Evaluation

8.1.1 Validating Implementation

Base Embeddings

As stated earlier, no publicly available version of the CoVeR algorithm exists. Being an extension of the Glove algorithm which still learns base word embeddings for focal and context vectors, comparing the outputs of both base embeddings served as a good metric. To measure the similarity between the two base embeddings generated by Glove and Cover, the F1 scores were generated to compare the

Covariate Specific Embeddings

The original CoVeR paper validates the quality of the learned covariate weight matrices by clustering. This was done

8.2 Model Evaluations

8.2.1 Language Model

8.2.2 Text Classification

8.3 SONGIFAI

8.3.1 Requirements Evaluation

8.3.2 Expert User Testing

Chapter 9

Conclusion

I was right all along.

9.1 What was I right about?

I was right about the following things.

9.1.1 Previous theories were wrong

People thought they understood, but they didn't.

9.1.2 My new idea is right

Of course.

Bibliography

- Bengio, Y., Ducharme, R., Vincent, P., and Janvin, C. (2003). A neural probabilistic language model. *J. Mach. Learn. Res.*, 3:1137–1155. [12](#)
- Chun Pui Tang, Ka Long Chui, Y. K. Y. Z. Z. K. H. W. (2018). Music genre classification using a hierarchical long short term memory (lstm) model. [8](#)
- Corinna Cortes, V. V. (1995). Support-vector networks. [15](#)
- Dai, A. M. and Le, Q. V. (2015). Semi-supervised sequence learning. [16](#)
- Edwards, P. and Rap, K. (2009). *How to Rap*. Chicago Review Press, Incorporated. [8](#)
- Harris, Z. S. (1954). Distributional structure. *Word*, 10(2-3):146–162. [1](#)
- Hochreiter, S. and Schmidhuber, J. (1997). Long short-term memory. *Neural computation*, 9(8):1735–1780. [14](#)
- Hulzebosch (2017). Deeplyricist: Automatic generation of rap lyrics using sequence-to-sequence learning. [8](#)
- Irvin, J., Chartock, E., and Hollander, N. (2016). Recurrent neural networks with attention for genre classification. [8](#)
- Joachims, T. (1998). Text categorization with support vector machines: Learning with many relevant features. In *European conference on machine learning*, pages 137–142. Springer. [15](#), [16](#)
- Lewis, D. D., Yang, Y., Rose, T. G., and Li, F. (2004). Rcv1: A new benchmark collection for text categorization research. [16](#)
- McCrum, R. (2011). *The Story of English*. BBC books. Faber & Faber. [2](#)
- Mikolov, T., Chen, K., Corrado, G., and Dean, J. (2013a). Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*. [1](#)
- Mikolov, T., Karafiát, M., Burget, L., Černocký, J., and Khudanpur, S. (2010). Recurrent neural network based language model. In *Eleventh annual conference of the international speech communication association*. [12](#)
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., and Dean, J. (2013b). Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, pages 3111–3119. [23](#)
- Pennington, J., Socher, R., and Manning, C. (2014). Glove: Global vectors for word representation. In *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)*, pages 1532–1543. [1](#), [10](#)

- Potash, P., Romanov, A., and Rumshisky, A. (2015). Ghostwriter: Using an lstm for automatic rap lyric generation. pages 1919–1924. [8](#)
- Rumelhart, D. E., Hinton, G. E., Williams, R. J., et al. (1986). Learning representations by back-propagating errors. *Cognitive modeling*. [1](#)
- Socher, R., Perelygin, A., Wu, J., Chuang, J., Manning, C. D., Ng, A., and Potts, C. (2013). Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pages 1631–1642. [9](#)
- Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2014). Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15:1929–1958. [24](#)
- Tian, K., Zhang, T., and Zou, J. (2018). Cover: Learning covariate-specific vector representations with tensor decompositions. *arXiv preprint arXiv:1802.07839*. [2](#)
- Zhang, X. and Lapata, M. (2014). Chinese poetry generation with recurrent neural networks. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 670–680. [8](#)

Appendix A

Code

```
10 PRINT "HELLO WORLD"
```